

---

# Assignment 1: Elliptic curve Diffie-Hellman (X25519)

Muhammad Ayain Fida Rana

Department of Computer Science and Technology

University of Cambridge

15 JJ Thomson Avenue, Cambridge CB3 0FD, United Kingdom

maf2@cam.ac.uk

## 1 Introduction

In this assignment, I have implemented the X25519 Diffie-Hellman key exchange, based on the Curve25519 elliptic curve, from scratch. My implementation supports two scalar multiplication algorithms: 1) Montgomery curve group law with a double-and-add algorithm, and 2) the Montgomery ladder. My implementation closely follows RFC 7748 ([Langley et al., 2016](#)) and Martin's tutorial ([Kleppmann, 2020](#)). The code is publicly available at <https://github.com/ayainfida/p79-assignment1-x25519>.

## 2 Implementation Architecture

The core implementation is split across modular components, briefly introduced here, and details follow in Section 3.

**Finite Fields (field.py).** Provides arithmetic operations (addition, subtraction, multiplication, squaring, inversion, and square-root operations.) on the prime field  $\mathbb{F}_p$ .

**Group Laws (group\_law.py).** Implements point addition and point doubling for Curve25519.

**Encoding/Decoding (encoding.py).** Handles conversion between bytes and integers for both scalars and  $x$ -coordinates.

**Scalar Multiplication Methods (methods.py).** Provides two scalar multiplication algorithms: a montgomery ladder and double-and-add.

**X25519 (api.py).** Exposes a public X25519 interface to generate a public/private key pair and then compute a shared secret.

**Defaults (default.py).** Defines shared constants used throughout the implementation.

## 3 Implementation Design

### 3.1 Finite Field Operations

Curve25519 is defined by the elliptic curve (Figure 1):

$$y^2 = x^3 + Ax^2 + x \quad \text{over the field } \mathbb{F}_p, \quad p = 2^{255} - 19, \quad A = 486662. \quad (1)$$

I did not attempt to justify the choice of parameters apart from their role in RFC 7748 ([Langley et al., 2016](#)). Specifically, the size of  $p$  determines the difficulty of the discrete logarithm problem and the performance of the curve operations.




Figure 1: Geometric visualization of an elliptic curve over the real numbers.

I represented the field elements using Python's `int (mod p)` after each operation. My implementation supports basic field operations including addition, subtraction, multiplication, squaring, inversion<sup>1</sup>, and division<sup>2</sup>.

Moreover, the double-and-add algorithm (detailed in Section 3.3.1) also requires the  $y$ -coordinate. I could have passed a valid  $y$  as described by Martin in class, but to avoid having different X25519 interfaces for both algorithms, I chose to calculate  $y$ , which required me to solve the following equation:

$$y = \pm \sqrt{x^3 + Ax^2 + x} \quad (2)$$

To compute the square root, I then defined an `fsqrt` operation based on the exponentiation-based method applicable to primes  $p \equiv 5 \pmod{8}$ , as described in RFC 8032 (Josefsson & Liusvaara, 2017).

### 3.2 Representation and Encoding

I defined a `Point` dataclass that stores  $(x, y)$  coordinate on Curve25519. If only  $x$  is provided,  $y$  is computed using the method `calculate_y(x)`. For the case where both  $(x, y)$  are provided, it validates this by checking if the following equation holds:

$$y^2 \pmod{p} \equiv x^3 + Ax^2 + x \pmod{p} \quad (3)$$

Moreover, according to the RFC 7748 specification (Langley et al., 2016):

- All 32-byte scalars and coordinates are encoded/decoded using the little-endian convention.
- The MSB of the last byte must be cleared when decoding the  $x$ -coordinate.

```
b = b[:-1] + bytes([b[-1] & 0x7F]) # unsets the MSB
return decode_little_endian(b) % P
```

The modulo  $p$  allows non-canonical values between  $2^{255} - 19$  and  $2^{255} - 1$  to be valid field elements as well.

- The scalars are clamped: 1) clear the 3 LSBs of the first byte, 2) clear the MSB of the last byte, and 3) set the second MSB of the last byte.

---

<sup>1</sup>implemented using Fermat's little theorem as in the lecture slides  
<sup>2</sup>by multiplying with its inverse

---

```

k_arr = bytearray(k)
k_arr[0] &= 248
k_arr[31] &= 127
k_arr[31] |= 64

```

### 3.3 Scalar Multiplication

#### 3.3.1 Double-and-Add

This method implements scalar multiplication using group laws. Given a scalar  $k$  and a valid point  $P$ , the algorithm recursively computes  $kP$  by repeated point doubling and point addition based on the value of  $k$ .

```

def double_and_add(k: int, Pt: Point) -> Point | PointAtInfinity:
    if k == 1:
        return Pt
    # k is even: Pt is doubled
    elif k & 1 == 0:
        return point_doubling(double_and_add(k // 2, Pt))
    # k is odd: Pt is doubled and then Pt is added
    else:
        return point_addition(point_doubling(double_and_add((k - 1) // 2, Pt)), Pt)

```

One thing to note here is that this method requires the complete point  $(x, y)$ . This is because point addition on an elliptic curve is geometrically defined (Figure 2) considering the straight line that passes through the two points<sup>3</sup> that are added.




Figure 2: Geometric illustration of elliptic curve point addition over the real numbers.

This line has a slope  $\lambda$ , which depends on both  $x$  and  $y$  and is used to determine the resulting point  $(x_3, y_3)$ , as shown in the equations<sup>4</sup> below:

$$x_3 = \lambda^2 - A - x_1 - x_2 \quad (4)$$

$$y_3 = \lambda(x_1 - x_3) - y_1 \quad (5)$$

#### 3.3.2 Montgomery Ladder

This method implements scalar multiplication using only  $x$ -coordinates. I implemented it following the code given in RFC 7748 ([Langley et al., 2016](#)) and Martin's tutorial ([Kleppmann](#),

---


<sup>3</sup>distinct

<sup>4</sup>following from Martin's tutorial ([Kleppmann, 2020](#))

2020). The idea is that it maintains two points that are a step apart and updates them in a fixed<sup>5</sup> pattern, unlike data-dependent sequences in double-and-add. In each iteration, the current scalar bit determines which of the two points represents the correct intermediate result. Moreover, projective coordinates of the form  $(X : Z) \equiv X/Z$  are used internally to delay field inversions<sup>6</sup> until the end, and are performed only once:  $\text{fdiv}(x\_2, z\_2)$ .

### 3.4 X25519 Key Exchange

The implemented X25519 interface allows setting the scalar multiplication algorithm to either LADDER or DOUBLE\_AND\_ADD. Each party generates a private key and derives the corresponding public key. They then compute a shared secret by applying X25519 to their private key and the peer's public key, which can then be hashed, e.g., using SHA256, to obtain a symmetric key.



```

# The code snippet demonstrates the relevant API calls.
x25519_instance = X25519(LADDER)

# Alice generates her key pair
alice_sk = x25519_instance.generate_private_key()
alice_pk = x25519_instance.derive_public_key(alice_sk)

# Bob generates his key pair
bob_sk = x25519_instance.generate_private_key()
bob_pk = x25519_instance.derive_public_key(bob_sk)

# Shared secret computation
alice_shared_secret = x25519_instance.x25519(alice_sk, bob_pk)
bob_shared_secret = x25519_instance.x25519(bob_sk, alice_pk)

# Derive a symmetric key
symmetric_key = sha256(alice_shared_secret).digest()
    
```

Figure 3: Elliptic Curve Diffie-Hellman Key Exchange

## 4 Testing and Validation

Individual tests are documented in `tests/*.py`; here, I provide a high-level overview of the testing strategy.

<sup>5</sup>18 arithmetic operations

<sup>6</sup>costly in terms of performance

---

**X25519 API Testing.** Validated my implementation against the official test vectors published in RFC 7748 ([Langley et al., 2016](#)), covering both single-shot and iterated scalar multiplication tests, and verified correct handling of invalid input lengths.

**Scalar Multiplication Agreement.** Verified that both scalar multiplication methods produce identical results for the same inputs.

**Diffie–Hellman (DH) Key Exchange and Agreement.** Tested DH key exchange with independent key pairs for both parties and verified that they derive the same shared secret. The key pairs were sourced from: a) RFC 7748 ([Langley et al., 2016](#)), b) `pycurve25519` ([TomCrypto, 2013](#)), and c) randomly generated private keys.

**Encoding.** Tested little-endian decoding/encoding,  $x$ -coordinate MSB masking (RFC 7748), and scalar clamping.

**Finite Field Operations** Tested for correctness of the defined operations and verified algebraic properties (commutativity, associativity, identity, inverse) for addition and multiplication.

**Group Law Operations** Validated the results of point addition and doubling are valid points on curve, along with edge cases like the point at infinity.

## 5 Discussion

### 5.1 Limitations and Production Considerations

I used Python’s built-in `int` type for field arithmetic over  $\mathbb{F}_p$ , which is not constant-time. This means that operations on values close to  $p$  may take longer than on smaller values, potentially leaking information through timing-based side-channel attacks. Therefore, my Montgomery ladder, despite its fixed pattern of operations, is vulnerable to such attacks. For a production-quality implementation, constant-time field arithmetic, uniform memory accesses, and avoidance of data-dependent branches<sup>7</sup> would be required.

Moreover, the double-and-add scalar multiplication method requires a valid  $(x, y)$  coordinate. While this holds for base point multiplication, it may fail for points that do not lie on the curve. I implemented `fsqrt` to compute the  $y$ -coordinate, but this is not tested except by following the methodology described in RFC 8032 ([Josefsson & Liusvaara, 2017](#)). Furthermore, I didn’t assess my implementation against active adversary attacks, which is a major concern in real-world scenarios.

### 5.2 Observations and Uncertainties

During testing, I observed that the second single-shot test vector from RFC 7748 did not succeed when using the double-and-add algorithm, resulting in an exception indicating that the given  $x$ -coordinate is not a valid point on Curve25519 (Figure 5.2). Upon validating with a reference<sup>8</sup> implementation, I found that after clearing the MSB of  $x$ , the resulting  $x$ -coordinate does not correspond to a valid point on Curve25519. In contrast, the Montgomery ladder successfully computes the correct result, highlighting its robustness against malformed public inputs.

Moreover, I was initially unsure whether to use standard integer division instead of field division when halving the scalar during recursion in the double-and-add algorithm. Since the scalar is clamped, it is supposed to be smaller than  $p$ , so I ended up using integer division.

---

<sup>7</sup>as in the double-and-add algorithm

<sup>8</sup><https://x25519.xargs.org/>



```
x-coordinate:  
    e5210f12786811d3f4b7959d0538ae2c31dbe7106fc03c3efc4cd549c715a493  
last byte:  
    0x93 -----> 0x13 (MSB cleared)  
updated x-coordinate:  
    e5210f12786811d3f4b7959d0538ae2c31dbe7106fc03c3efc4cd549c715a413  
big-endian format for tool:  
    13a415c749d54cf3e3cc06f10e7db312cae38059d95b7f4d3116878120f21e5
```

Figure 4: Absence of a valid curve point for the RFC 7748 test vector 2.

## References

- Simon Josefsson and Ilari Liusvaara. Edwards-Curve Digital Signature Algorithm (EdDSA). RFC 8032, January 2017. URL <https://www.rfc-editor.org/info/rfc8032>.
- Martin Kleppmann. Implementing curve25519/x25519: A tutorial on elliptic curve cryptography. 2020. URL <https://martin.kleppmann.com/papers/curve25519.pdf>.
- Adam Langley, Mike Hamburg, and Sean Turner. Elliptic Curves for Security. RFC 7748, January 2016. URL <https://www.rfc-editor.org/info/rfc7748>.
- TomCrypto. pycurve25519. <https://github.com/TomCrypto/pycurve25519>, 2013.